

# Different MLP Arcitectures on MNIST dataset

## Loading data

```
In [2]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow"
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
In [0]: import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid(True)
    fig.canvas.draw()
```

```
In [4]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz> (<https://s3.amazonaws.com/img-datasets/mnist.npz>)  
 11493376/11490434 [=====] - 1s 0us/step

```
In [5]: print("Number of training examples :", X_train.shape[0], "and each image is of shape", X_train.shape[1]*X_train.shape[2])
print("Number of training examples :", X_test.shape[0], "and each image is of shape", X_test.shape[1]*X_test.shape[2])
```

Number of training examples : 60000 and each image is of shape (28, 28)  
 Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [7]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (784)")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (784)")
```

Number of training examples : 60000 and each image is of shape (784)  
 Number of training examples : 10000 and each image is of shape (784)

```
In [8]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18  126  136  175  26  166  255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94  154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
In [10]: # example data point after normalizing
print(X_train[0])
```

```
[0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.]
```

```
In [11]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

```
In [0]: from keras.models import Sequential
from keras.layers import Dense, Activation

# some model parameters
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

# [1] MLP with 2 hidden layers

## [ 1.1] Without Batch Normalization and Dropout

```
In [15]: # Multilayer perceptron

model = Sequential()
model.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_4 (Dense)	(None, 364)	285740
dense_5 (Dense)	(None, 52)	18980
dense_6 (Dense)	(None, 10)	530
=====	=====	=====
Total params: 305,250		
Trainable params: 305,250		
Non-trainable params: 0		

```
In [16]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 86us/step - loss: 0.2741 - acc: 0.9213 - val_loss: 0.1435 - val_acc: 0.9563
Epoch 2/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1041 - acc: 0.9690 - val_loss: 0.0936 - val_acc: 0.9729
Epoch 3/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0665 - acc: 0.9796 - val_loss: 0.0862 - val_acc: 0.9755
Epoch 4/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0468 - acc: 0.9855 - val_loss: 0.0678 - val_acc: 0.9787
Epoch 5/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0343 - acc: 0.9892 - val_loss: 0.0663 - val_acc: 0.9803
Epoch 6/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0247 - acc: 0.9924 - val_loss: 0.0746 - val_acc: 0.9777
Epoch 7/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0221 - acc: 0.9933 - val_loss: 0.0642 - val_acc: 0.9820
Epoch 8/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0161 - acc: 0.9946 - val_loss: 0.0775 - val_acc: 0.9788
Epoch 9/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0134 - acc: 0.9957 - val_loss: 0.0802 - val_acc: 0.9788
Epoch 10/20
60000/60000 [=====] - 5s 75us/step - loss: 0.0113 - acc: 0.9964 - val_loss: 0.0733 - val_acc: 0.9803
Epoch 11/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0126 - acc: 0.9960 - val_loss: 0.0737 - val_acc: 0.9807
Epoch 12/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0090 - acc: 0.9973 - val_loss: 0.0772 - val_acc: 0.9816
Epoch 13/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0088 - acc: 0.9971 - val_loss: 0.0790 - val_acc: 0.9812
Epoch 14/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0079 - acc: 0.9975 - val_loss: 0.0877 - val_acc: 0.9808
Epoch 15/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0071 - acc: 0.9976 - val_loss: 0.0893 - val_acc: 0.9781
Epoch 16/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0113 - acc: 0.9965 - val_loss: 0.0975 - val_acc: 0.9787
Epoch 17/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0045 - acc: 0.9986 - val_loss: 0.0982 - val_acc: 0.9792
Epoch 18/20
```

```

60000/60000 [=====] - 4s 73us/step - loss: 0.0110 - ac
c: 0.9965 - val_loss: 0.0809 - val_acc: 0.9820
Epoch 19/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0037 - ac
c: 0.9990 - val_loss: 0.0829 - val_acc: 0.9831
Epoch 20/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0056 - ac
c: 0.9982 - val_loss: 0.0965 - val_acc: 0.9793

```

```

In [18]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09650619211766216

Test accuracy: 0.9793



```

In [19]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure(figsize=(8, 6))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.  
 kde\_data = remove\_na(group\_data)  
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.  
 violin\_data = remove\_na(group\_data)



In [0]:

## [1.2] With Batch Normalization and Dropout

```
In [21]: from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model = Sequential()
model.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0,
model.add(BatchNormalization()))
model.add(Dropout(0.5))
model.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(BatchNormalization()))
model.add(Dropout(0.5))
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 364)	285740
batch_normalization_1 (Batch Normalization)	(None, 364)	1456
dropout_1 (Dropout)	(None, 364)	0
dense_8 (Dense)	(None, 52)	18980
batch_normalization_2 (Batch Normalization)	(None, 52)	208
dropout_2 (Dropout)	(None, 52)	0
dense_9 (Dense)	(None, 10)	530
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		



```
In [22]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 107us/step - loss: 0.5517 - accuracy: 0.8330 - val_loss: 0.1736 - val_acc: 0.9462
Epoch 2/20
60000/60000 [=====] - 6s 96us/step - loss: 0.2584 - accuracy: 0.9258 - val_loss: 0.1278 - val_acc: 0.9600
Epoch 3/20
60000/60000 [=====] - 6s 95us/step - loss: 0.2003 - accuracy: 0.9423 - val_loss: 0.1044 - val_acc: 0.9666
Epoch 4/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1691 - accuracy: 0.9502 - val_loss: 0.0951 - val_acc: 0.9706
Epoch 5/20
60000/60000 [=====] - 6s 96us/step - loss: 0.1487 - accuracy: 0.9580 - val_loss: 0.0830 - val_acc: 0.9750
Epoch 6/20
60000/60000 [=====] - 6s 96us/step - loss: 0.1339 - accuracy: 0.9607 - val_loss: 0.0770 - val_acc: 0.9758
Epoch 7/20
60000/60000 [=====] - 6s 96us/step - loss: 0.1257 - accuracy: 0.9637 - val_loss: 0.0756 - val_acc: 0.9786
Epoch 8/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1200 - accuracy: 0.9646 - val_loss: 0.0714 - val_acc: 0.9774
Epoch 9/20
60000/60000 [=====] - 6s 97us/step - loss: 0.1103 - accuracy: 0.9667 - val_loss: 0.0714 - val_acc: 0.9779
Epoch 10/20
60000/60000 [=====] - 6s 98us/step - loss: 0.1035 - accuracy: 0.9689 - val_loss: 0.0715 - val_acc: 0.9781
Epoch 11/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0993 - accuracy: 0.9704 - val_loss: 0.0673 - val_acc: 0.9794
Epoch 12/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0945 - accuracy: 0.9718 - val_loss: 0.0667 - val_acc: 0.9806
Epoch 13/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0905 - accuracy: 0.9728 - val_loss: 0.0634 - val_acc: 0.9803
Epoch 14/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0890 - accuracy: 0.9733 - val_loss: 0.0647 - val_acc: 0.9805
Epoch 15/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0835 - accuracy: 0.9753 - val_loss: 0.0686 - val_acc: 0.9796
Epoch 16/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0787 - accuracy: 0.9765 - val_loss: 0.0614 - val_acc: 0.9815
Epoch 17/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0801 - accuracy: 0.9764 - val_loss: 0.0632 - val_acc: 0.9816
Epoch 18/20
```

```

60000/60000 [=====] - 6s 95us/step - loss: 0.0736 - ac
c: 0.9782 - val_loss: 0.0661 - val_acc: 0.9808
Epoch 19/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0753 - ac
c: 0.9775 - val_loss: 0.0642 - val_acc: 0.9819
Epoch 20/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0719 - ac
c: 0.9784 - val_loss: 0.0645 - val_acc: 0.9819

```

```

In [23]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06447895890461514

Test accuracy: 0.9819



```

In [24]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(8, 6))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

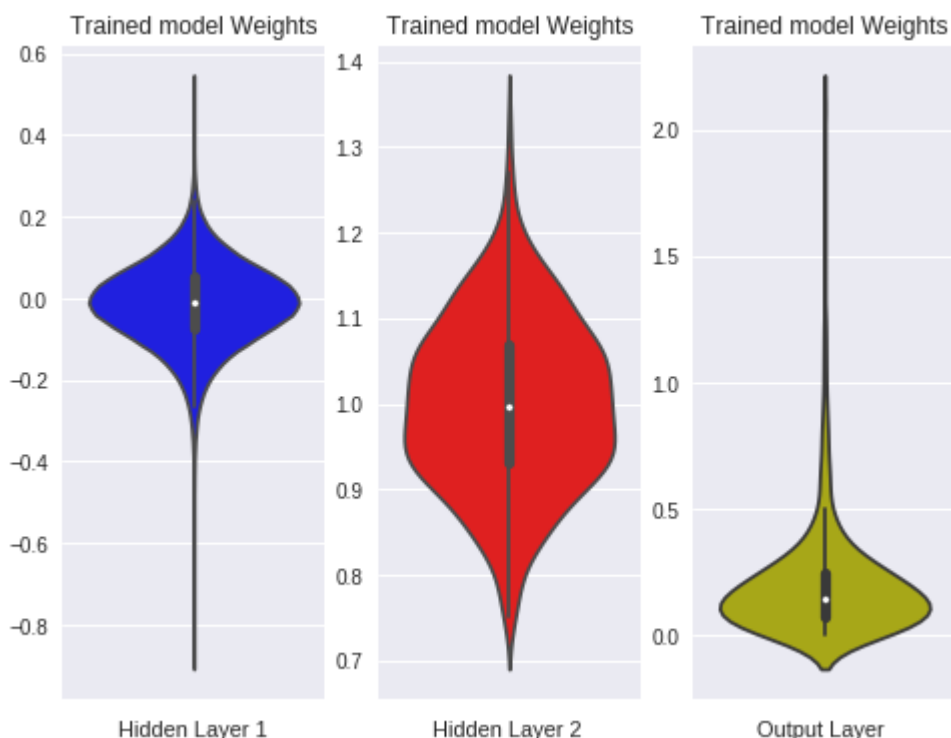
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

kde\_data = remove\_na(group\_data)

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

violin\_data = remove\_na(group\_data)



## [2] MLP with 3 hidden layers

### [2.1] Without Batch Normalization and Dropout

```
In [25]: # Multilayer perceptron

model = Sequential()
model.add(Dense(392, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(196, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 392)	307720
dense_11 (Dense)	(None, 196)	77028
dense_12 (Dense)	(None, 98)	19306
dense_13 (Dense)	(None, 10)	990
Total params: 405,044		
Trainable params: 405,044		
Non-trainable params: 0		

```
In [26]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 102us/step - loss: 0.2474 - acc: 0.9268 - val_loss: 0.1271 - val_acc: 0.9592
Epoch 2/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0884 - acc: 0.9727 - val_loss: 0.0834 - val_acc: 0.9732
Epoch 3/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0570 - acc: 0.9820 - val_loss: 0.0759 - val_acc: 0.9760
Epoch 4/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0415 - acc: 0.9869 - val_loss: 0.0653 - val_acc: 0.9793
Epoch 5/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0305 - acc: 0.9900 - val_loss: 0.0713 - val_acc: 0.9800
Epoch 6/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0247 - acc: 0.9919 - val_loss: 0.0776 - val_acc: 0.9788
Epoch 7/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0207 - acc: 0.9929 - val_loss: 0.0804 - val_acc: 0.9792
Epoch 8/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0184 - acc: 0.9938 - val_loss: 0.0773 - val_acc: 0.9804
Epoch 9/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0194 - acc: 0.9934 - val_loss: 0.0800 - val_acc: 0.9778
Epoch 10/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0143 - acc: 0.9951 - val_loss: 0.0821 - val_acc: 0.9789
Epoch 11/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0133 - acc: 0.9956 - val_loss: 0.0781 - val_acc: 0.9820
Epoch 12/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0141 - acc: 0.9951 - val_loss: 0.0954 - val_acc: 0.9775
Epoch 13/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0108 - acc: 0.9967 - val_loss: 0.0913 - val_acc: 0.9775
Epoch 14/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0141 - acc: 0.9954 - val_loss: 0.0915 - val_acc: 0.9779
Epoch 15/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0096 - acc: 0.9968 - val_loss: 0.0893 - val_acc: 0.9795
Epoch 16/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0101 - acc: 0.9970 - val_loss: 0.0920 - val_acc: 0.9800
Epoch 17/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0116 - acc: 0.9965 - val_loss: 0.0950 - val_acc: 0.9808
Epoch 18/20
```

```

60000/60000 [=====] - 6s 96us/step - loss: 0.0089 - ac
c: 0.9972 - val_loss: 0.0992 - val_acc: 0.9789
Epoch 19/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0073 - ac
c: 0.9976 - val_loss: 0.0874 - val_acc: 0.9823
Epoch 20/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0057 - ac
c: 0.9982 - val_loss: 0.0892 - val_acc: 0.9828

```

```

In [27]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08915272126807804

Test accuracy: 0.9828



```

In [28]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(8, 6))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()


```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

kde\_data = remove\_na(group\_data)

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

violin\_data = remove\_na(group\_data)



## [2.2] With Batch Normalization and Dropout

```
In [29]: from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model = Sequential()
model.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 364)	285740
batch_normalization_3 (Batch Normalization)	(None, 364)	1456
dropout_3 (Dropout)	(None, 364)	0
dense_15 (Dense)	(None, 52)	18980
batch_normalization_4 (Batch Normalization)	(None, 52)	208
dropout_4 (Dropout)	(None, 52)	0
dense_16 (Dense)	(None, 64)	3392
batch_normalization_5 (Batch Normalization)	(None, 64)	256
dropout_5 (Dropout)	(None, 64)	0
dense_17 (Dense)	(None, 10)	650
Total params: 310,682		
Trainable params: 309,722		
Non-trainable params: 960		



```
In [31]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 127us/step - loss: 0.8095 - accuracy: 0.7476 - val_loss: 0.2088 - val_acc: 0.9344
Epoch 2/20
60000/60000 [=====] - 6s 107us/step - loss: 0.3403 - accuracy: 0.9033 - val_loss: 0.1457 - val_acc: 0.9562
Epoch 3/20
60000/60000 [=====] - 6s 106us/step - loss: 0.2610 - accuracy: 0.9285 - val_loss: 0.1263 - val_acc: 0.9634
Epoch 4/20
60000/60000 [=====] - 6s 105us/step - loss: 0.2190 - accuracy: 0.9390 - val_loss: 0.1025 - val_acc: 0.9711
Epoch 5/20
60000/60000 [=====] - 6s 104us/step - loss: 0.1953 - accuracy: 0.9472 - val_loss: 0.0967 - val_acc: 0.9721
Epoch 6/20
60000/60000 [=====] - 6s 105us/step - loss: 0.1848 - accuracy: 0.9492 - val_loss: 0.0968 - val_acc: 0.9724
Epoch 7/20
60000/60000 [=====] - 6s 104us/step - loss: 0.1687 - accuracy: 0.9546 - val_loss: 0.0914 - val_acc: 0.9737
Epoch 8/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1545 - accuracy: 0.9578 - val_loss: 0.0838 - val_acc: 0.9762
Epoch 9/20
60000/60000 [=====] - 6s 104us/step - loss: 0.1435 - accuracy: 0.9606 - val_loss: 0.0879 - val_acc: 0.9756
Epoch 10/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1431 - accuracy: 0.9607 - val_loss: 0.0794 - val_acc: 0.9787
Epoch 11/20
60000/60000 [=====] - 7s 109us/step - loss: 0.1350 - accuracy: 0.9622 - val_loss: 0.0813 - val_acc: 0.9769
Epoch 12/20
60000/60000 [=====] - 6s 107us/step - loss: 0.1268 - accuracy: 0.9652 - val_loss: 0.0741 - val_acc: 0.9799
Epoch 13/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1232 - accuracy: 0.9660 - val_loss: 0.0773 - val_acc: 0.9792
Epoch 14/20
60000/60000 [=====] - 6s 105us/step - loss: 0.1180 - accuracy: 0.9678 - val_loss: 0.0735 - val_acc: 0.9793
Epoch 15/20
60000/60000 [=====] - 6s 107us/step - loss: 0.1122 - accuracy: 0.9694 - val_loss: 0.0723 - val_acc: 0.9808
Epoch 16/20
60000/60000 [=====] - 6s 108us/step - loss: 0.1082 - accuracy: 0.9704 - val_loss: 0.0732 - val_acc: 0.9804
Epoch 17/20
60000/60000 [=====] - 6s 107us/step - loss: 0.1081 - accuracy: 0.9701 - val_loss: 0.0731 - val_acc: 0.9805
Epoch 18/20
```

```

60000/60000 [=====] - 7s 108us/step - loss: 0.1008 - a
cc: 0.9727 - val_loss: 0.0738 - val_acc: 0.9801
Epoch 19/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1008 - a
cc: 0.9720 - val_loss: 0.0692 - val_acc: 0.9821
Epoch 20/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0989 - a
cc: 0.9724 - val_loss: 0.0764 - val_acc: 0.9795

```

```

In [32]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07640664648028324

Test accuracy: 0.9795



```

In [33]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(8, 6))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

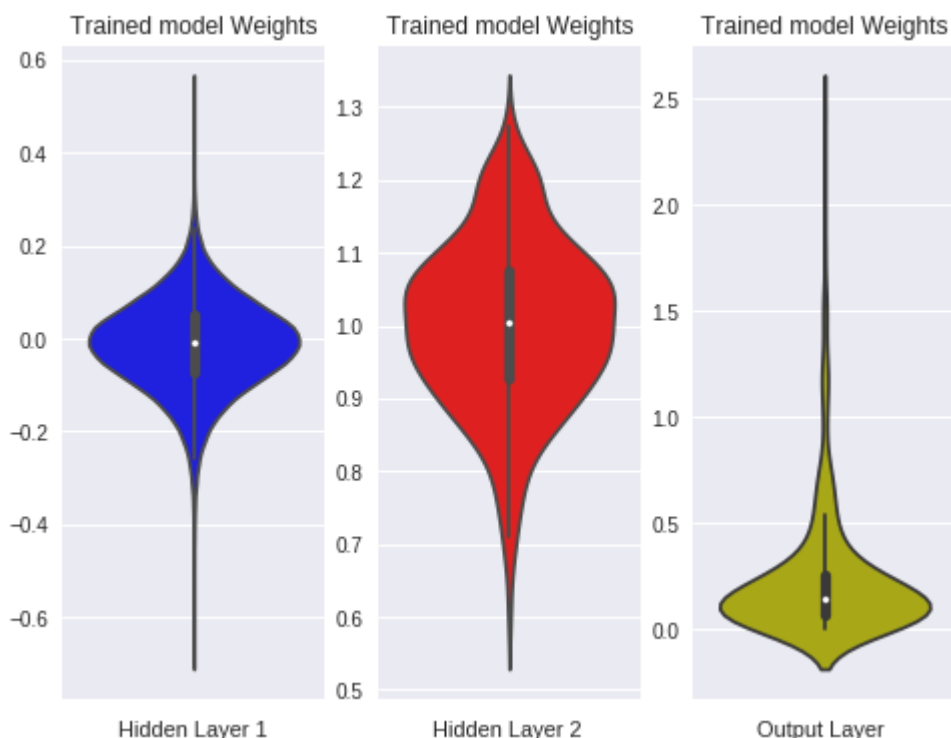
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

kde\_data = remove\_na(group\_data)

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

violin\_data = remove\_na(group\_data)



## [2] MLP with 5 hidden layers

### [3.1] Without Batch Normalization and Dropout

```
In [34]: # Multilayer perceptron

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_18 (Dense)	(None, 512)	401920
dense_19 (Dense)	(None, 256)	131328
dense_20 (Dense)	(None, 128)	32896
dense_21 (Dense)	(None, 64)	8256
dense_22 (Dense)	(None, 32)	2080
dense_23 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		
=====		

```
In [35]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 9s 142us/step - loss: 0.3075 - accuracy: 0.9062 - val_loss: 0.1160 - val_acc: 0.9641
Epoch 2/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0998 - accuracy: 0.9698 - val_loss: 0.1045 - val_acc: 0.9681
Epoch 3/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0650 - accuracy: 0.9800 - val_loss: 0.0769 - val_acc: 0.9760
Epoch 4/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0453 - accuracy: 0.9860 - val_loss: 0.0769 - val_acc: 0.9772
Epoch 5/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0388 - accuracy: 0.9875 - val_loss: 0.0905 - val_acc: 0.9762
Epoch 6/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0311 - accuracy: 0.9901 - val_loss: 0.0786 - val_acc: 0.9794
Epoch 7/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0260 - accuracy: 0.9920 - val_loss: 0.0770 - val_acc: 0.9793
Epoch 8/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0215 - accuracy: 0.9934 - val_loss: 0.0936 - val_acc: 0.9768
Epoch 9/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0196 - accuracy: 0.9937 - val_loss: 0.0951 - val_acc: 0.9768
Epoch 10/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0166 - accuracy: 0.9945 - val_loss: 0.0862 - val_acc: 0.9799
Epoch 11/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0169 - accuracy: 0.9948 - val_loss: 0.0959 - val_acc: 0.9774
Epoch 12/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0166 - accuracy: 0.9948 - val_loss: 0.0949 - val_acc: 0.9774
Epoch 13/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0137 - accuracy: 0.9956 - val_loss: 0.0884 - val_acc: 0.9803
Epoch 14/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0127 - accuracy: 0.9960 - val_loss: 0.0816 - val_acc: 0.9808
Epoch 15/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0139 - accuracy: 0.9958 - val_loss: 0.1052 - val_acc: 0.9788
Epoch 16/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0130 - accuracy: 0.9963 - val_loss: 0.0854 - val_acc: 0.9832
Epoch 17/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0104 - accuracy: 0.9971 - val_loss: 0.0868 - val_acc: 0.9818
Epoch 18/20
```

```

60000/60000 [=====] - 8s 133us/step - loss: 0.0126 - a
cc: 0.9963 - val_loss: 0.0794 - val_acc: 0.9822
Epoch 19/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0097 - a
cc: 0.9973 - val_loss: 0.1044 - val_acc: 0.9797
Epoch 20/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0110 - a
cc: 0.9968 - val_loss: 0.1060 - val_acc: 0.9785

```

```

In [36]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10602583270173854

Test accuracy: 0.9785



```

In [37]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(8, 6))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()


```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

kde\_data = remove\_na(group\_data)

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

violin\_data = remove\_na(group\_data)



### [3.2] With Batch Normalization and Dropout

```
In [38]: from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(output_dim, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 512)	401920
batch_normalization_6 (Batch Normalization)	(None, 512)	2048
dropout_6 (Dropout)	(None, 512)	0
dense_25 (Dense)	(None, 256)	131328
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dropout_7 (Dropout)	(None, 256)	0
dense_26 (Dense)	(None, 128)	32896
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dropout_8 (Dropout)	(None, 128)	0
dense_27 (Dense)	(None, 64)	8256
batch_normalization_9 (Batch Normalization)	(None, 64)	256
dropout_9 (Dropout)	(None, 64)	0
dense_28 (Dense)	(None, 32)	2080
batch_normalization_10 (Batch Normalization)	(None, 32)	128
dropout_10 (Dropout)	(None, 32)	0



dense_29 (Dense)	(None, 10)	330
------------------	------------	-----

---

Total params: 580,778  
Trainable params: 578,794  
Non-trainable params: 1,984

---

```
In [39]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ver
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 13s 224us/step - loss: 1.2762 - acc: 0.5948 - val\_loss: 0.2912 - val\_acc: 0.9224

Epoch 2/20

60000/60000 [=====] - 11s 190us/step - loss: 0.5094 - acc: 0.8618 - val\_loss: 0.1788 - val\_acc: 0.9523

Epoch 3/20

60000/60000 [=====] - 11s 190us/step - loss: 0.3592 - acc: 0.9095 - val\_loss: 0.1466 - val\_acc: 0.9610

Epoch 4/20

60000/60000 [=====] - 11s 188us/step - loss: 0.2985 - acc: 0.9285 - val\_loss: 0.1363 - val\_acc: 0.9660

Epoch 5/20

60000/60000 [=====] - 11s 188us/step - loss: 0.2592 - acc: 0.9387 - val\_loss: 0.1195 - val\_acc: 0.9697

Epoch 6/20

60000/60000 [=====] - 11s 190us/step - loss: 0.2315 - acc: 0.9456 - val\_loss: 0.1222 - val\_acc: 0.9695

Epoch 7/20

60000/60000 [=====] - 11s 190us/step - loss: 0.2225 - acc: 0.9487 - val\_loss: 0.1162 - val\_acc: 0.9709

Epoch 8/20

60000/60000 [=====] - 12s 192us/step - loss: 0.2027 - acc: 0.9530 - val\_loss: 0.0990 - val\_acc: 0.9757

Epoch 9/20

60000/60000 [=====] - 12s 198us/step - loss: 0.1941 - acc: 0.9557 - val\_loss: 0.0959 - val\_acc: 0.9758

Epoch 10/20

60000/60000 [=====] - 11s 190us/step - loss: 0.1793 - acc: 0.9581 - val\_loss: 0.0988 - val\_acc: 0.9770

Epoch 11/20

60000/60000 [=====] - 11s 189us/step - loss: 0.1725 - acc: 0.9602 - val\_loss: 0.0918 - val\_acc: 0.9790

Epoch 12/20

60000/60000 [=====] - 11s 189us/step - loss: 0.1606 - acc: 0.9622 - val\_loss: 0.0887 - val\_acc: 0.9801

Epoch 13/20

60000/60000 [=====] - 11s 190us/step - loss: 0.1556 - acc: 0.9647 - val\_loss: 0.0910 - val\_acc: 0.9787

Epoch 14/20

60000/60000 [=====] - 11s 190us/step - loss: 0.1521 - acc: 0.9651 - val\_loss: 0.0865 - val\_acc: 0.9796

Epoch 15/20

60000/60000 [=====] - 11s 189us/step - loss: 0.1467 - acc: 0.9675 - val\_loss: 0.0810 - val\_acc: 0.9808

Epoch 16/20

60000/60000 [=====] - 11s 189us/step - loss: 0.1424 - acc: 0.9677 - val\_loss: 0.0794 - val\_acc: 0.9807

Epoch 17/20

60000/60000 [=====] - 11s 187us/step - loss: 0.1334 - acc: 0.9695 - val\_loss: 0.0847 - val\_acc: 0.9816

Epoch 18/20

```

60000/60000 [=====] - 11s 190us/step - loss: 0.1328 -
acc: 0.9700 - val_loss: 0.0826 - val_acc: 0.9808
Epoch 19/20
60000/60000 [=====] - 11s 190us/step - loss: 0.1287 -
acc: 0.9709 - val_loss: 0.0780 - val_acc: 0.9822
Epoch 20/20
60000/60000 [=====] - 11s 191us/step - loss: 0.1271 -
acc: 0.9715 - val_loss: 0.0801 - val_acc: 0.9819

```

```

In [40]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08005840601597447

Test accuracy: 0.9819



```

In [41]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(8, 6))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

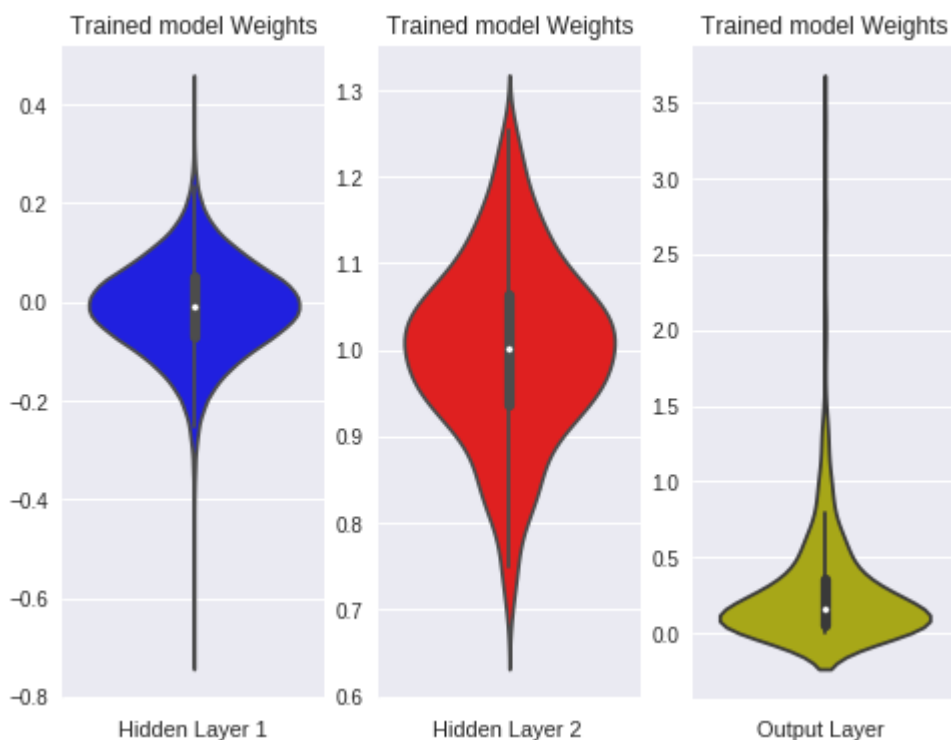
```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

kde\_data = remove\_na(group\_data)

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove\_na is deprecated and is a private function. Do not use.

violin\_data = remove\_na(group\_data)



```
In [47]: from prettytable import PrettyTable
x = PrettyTable()

x.field_names = ["MLP_MODEL", "TRAIN_ACCURACY", "TEST_ACCURACY"]
x.add_row(["MLP(2-hidden layers)", 0.870, 0.875])
x.add_row(["MLP(2-hidden layers) With Dropout and Batch Normalization", 0.978, 0.978])
x.add_row(["MLP(3-hidden layers)", 0.998, 0.982])
x.add_row(["MLP(3-hidden layers) With Dropout and Batch Normalization", 0.972, 0.972])
x.add_row(["MLP(5-hidden layers)", 0.996, 0.978])
x.add_row(["MLP(5-hidden layers) With Dropout and Batch Normalization", 0.971, 0.971])

print('\t\t\t\tMLP WITH DIFFERNET ARCHITECTURES')
print(x)
```

TEST_ACCURACY	MLP_MODEL	TRAIN_ACCURACY
0.875	MLP(2-hidden layers)	0.87
0.981	MLP(2-hidden layers) With Dropout and Batch Normalization	0.978
0.982	MLP(3-hidden layers)	0.998
0.979	MLP(3-hidden layers) With Dropout and Batch Normalization	0.972
0.978	MLP(5-hidden layers)	0.996
0.981	MLP(5-hidden layers) With Dropout and Batch Normalization	0.971

## Conclusion

1. The Best train and test (0.998 | 0.982) accuracy is obtained by MLP(3-hidden layers) without Dropout and Batch Normalization.
2. With Dropout and Batch Normalization the NN is less overfitted than the NN without Dropout and Batch Normalization.